

Bayesian analysis of the NESTA study of interventions against verbal aggression online

Rafal Urbaniak

Contents

1	Exploration	1
2	Causal inference	7
3	Bayesian models, priors and diagnostics	10
	References	15

1 Exploration

Load the dataset and take a look first.

```
summaries <- read.csv(file = "datasets/Summaries.csv")
head(summaries) %>% kable( "latex", booktabs = T) %>%
  kable_styling(latex_options = c("striped", "scale_down"), font_size = 9)
```

The basic variables we are dealing with are in the following table.

Further variables are defined in terms of those, in particular, we will be predicting AdiffS which is the standardized difference AA-AB, and AdiffC, which is the standardized difference CA-CB. Before we proceed, we will also standardize the predictors, and add a numerical index for the group:

```
summaries$ABS <- standardize(summaries$AB)
summaries$CBS <- standardize(summaries$CB)
summaries$AAS <- standardize(summaries$AA)
summaries$CAS <- standardize(summaries$CA)
summaries$CDS <- standardize(summaries$CD)
summaries$ADS <- standardize(summaries$AD)
summaries$group <- as.factor(summaries$group)
summaries$groupID <- as.integer( as.factor(summaries$group) )
```


First, let's take a look at the distribution of IC in the treatment groups:

```
ggplot(summaries[summaries$group != "control",], aes(x = IC, fill = group)) +
  geom_bar() + theme_tufte() +
  xlab("interventions received") +
  labs(title = "Intervention counts in treatment groups") +
  scale_x_continuous(breaks = seq(0, 40, 5))
```

X	author	AB	AD	AA	CB	CD	CA	Adiff	Cdiff	AdiffS	CdiffS	group	IC
1	_swf	19	1	0	720	25	28	-19	-692	-0.0245122	-0.3501491	normative	1
2	-Allergic	24	24	8	1614	1451	1237	-16	-377	0.0719197	0.1057675	normative	3
3	-funny-username-	23	6	12	847	497	721	-11	-126	0.2326395	0.4690535	control	0
4	-Johnny-	18	2	8	1465	408	684	-10	-781	0.2647835	-0.4789637	empathy	2
5	1secwhileiyet3	15	3	4	1384	198	120	-11	-1264	0.2326395	-1.1780359	control	0
6	20CharsIsNotEnough	16	10	25	779	907	972	9	193	0.8755188	0.9307596	empathy	4

variable	explanation
AB	attacks before (pre-treatment)
AD	attacks during (the treatment period)
AA	attacks after (post-treatment)
CB	comments before
CD	comments during
CA	comments after
group	treatment group
IC	intervention count

Intervention counts in treatment groups



Second, when we look at the distribution of standardized difference in attacks, when restricted to $(-1,1)$, the peaks of distributions are shifted a bit, with lowest median for the normative group, but not too much:

```
violAdiffS <- ggplot(summaries, aes(x=group, y = AdiffS)) +
  geom_violin() + theme_tufte()
violJoint <- ggarrange(violAdiffS+ggtitle("whole range"),
  violAdiffS + ylim(c(-1,1))+geom_boxplot(width = .2)+
  ggtitle("restricted to (-1,1)"))
```

```
## Warning: Removed 58 rows containing non-finite values (stat_ydensity).
```

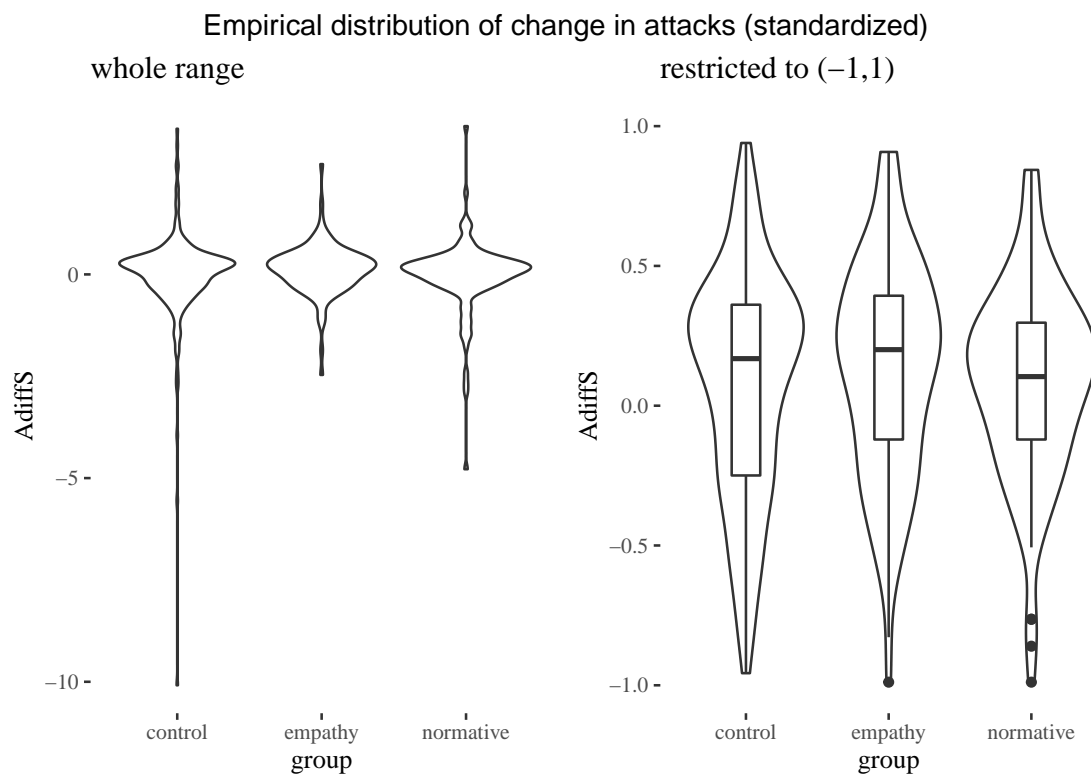
```
## Warning: Removed 58 rows containing non-finite values (stat_boxplot).
```

```
violJointTitled <- annotate_figure(violJoint,
  top = text_grob("Empirical distribution of change in attacks (standardized)",
    size = 12))
```

```
violJointTitled
```

Note there were much more empathetic interventions, this needs an explanation

Question: intervention counts by group




Analogous plot for comments does not reveal this slight downward shift for normative, but otherwise the visualisation might suggest no strong impact of interventions on attacks, and no impact on comments.

```
violCdiffs <- ggplot(summaries, aes(x=group, y = Cdiffs))+
  geom_violin() +theme_tufte()
violJointC <- ggarrange(violCdiffs+ggtitle("whole range"),
  violCdiffs + ylim(c(-1,1))+geom_boxplot(width = .2)+
  ggtitle("restricted to (-1,1)"))

## Warning: Removed 90 rows containing non-finite values (stat_ydensity).
## Warning: Removed 90 rows containing non-finite values (stat_boxplot).

violJointCTitled <- annotate_figure(violJoint,
  top = text_grob("Empirical distribution of change in comments (standardized)",
    size = 12))
violJointCTitled
```




However, plotting changes against intervention counts reveals that restricting attention to various activity levels drastically changes the regression lines.

```
icplot1 <- ggplot(summaries, aes(x = IC, y = AdiffS, color = group, fill = group)) +
  geom_jitter(alpha = 0.6, size = .8) + theme_tufte() +
  geom_smooth(alpha = 0.2, method = "lm") +
  xlim(c(0,25)) + ylim(c(-2,2)) +
  ggtitle("sd restricted to (-2,2)") +
  theme(legend.position = c(0.65, 0.1))

icplot2 <- ggplot(summaries, aes(x = IC, y = AdiffS, color = group, fill = group)) +
  geom_jitter(alpha = 0.6, size = .8) + theme_tufte() +
  geom_smooth(alpha = 0.2, method = "lm") +
  xlim(c(0,25)) + ylim(c(-1,1)) + ggtitle("sd restricted to (-1,1)") +
  theme(legend.position = c(0.65, 0.1))

icplotJoint <- ggarrange(icplot1, icplot2)
icplotTitled <- annotate_figure(icplotJoint,
  top = text_grob("Change in attacks (standardized) vs interventions received", size = 12))
icplotTitled
```

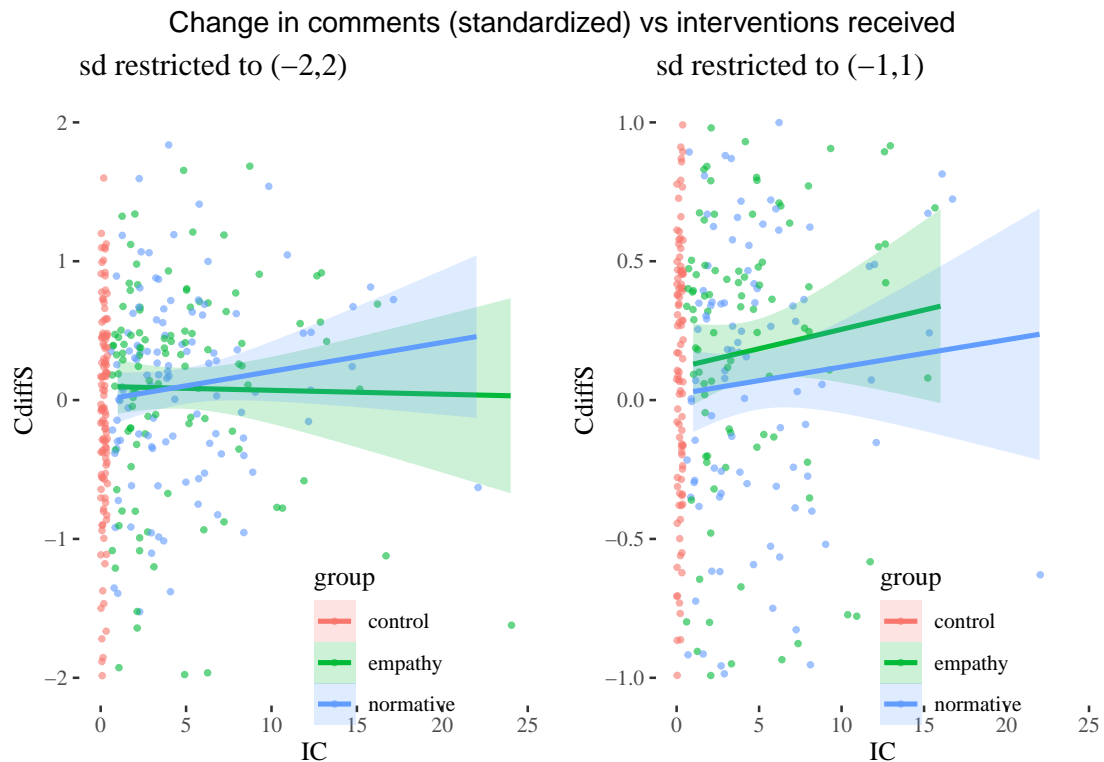


Some interactions are also suggested by the differences in linear smoothing when attention is restricted when it comes to change in comments.

```
icCplot1 <- ggplot(summaries, aes(x = IC, y = CdiffS, color = group, fill = group)) +
  geom_jitter(alpha = 0.6, size = .8) + theme_tufte() +
  geom_smooth(alpha = 0.2, method = "lm") +
  xlim(c(0,25)) + ylim(c(-2,2)) +
  ggtitle("sd restricted to (-2,2)") +
  theme(legend.position = c(0.65, 0.1))

icCplot2 <- ggplot(summaries, aes(x = IC, y = CdiffS, color = group, fill = group)) +
  geom_jitter(alpha = 0.6, size = .8) + theme_tufte() +
  geom_smooth(alpha = 0.2, method = "lm") +
  xlim(c(0,25)) + ylim(c(-1,1)) + ggtitle("sd restricted to (-1,1)") +
  theme(legend.position = c(0.65, 0.1))

icCplotJoint <- ggarrange(icCplot1, icCplot2)
icCplotTitled <- annotate_figure(icCplotJoint,
  top = text_grob("Change in comments (standardized) vs interventions received",
    size = 12))
icCplotTitled
```



This suggests we should keep an eye out for interactions in the analysis, and that the initial comparison of means or medians between groups might be misleading if the effects in different volume groups are different and cancel each other.

Now, let's inspect correlations between the variables involved in the model:

```
summariesCorr <- select(summaries, IC, ABS, CBS, AAS, CAS, CDS, ADS)
ggcorr(summariesCorr, method = c("pairwise"),
  digits = 4, low = "steelblue", mid = "white",
  high = "darkred", midpoint = 0,
  geom = "tile", label = TRUE, label_size = 4, label_round = 2, layout.exp = 1,
  label_alpha = FALSE, hjust = 0.75)
```



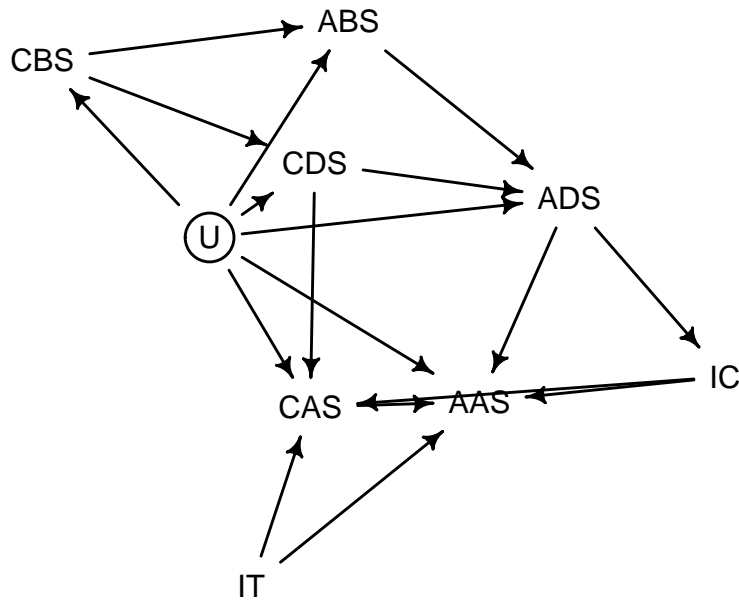
This tells us that almost no predictors are strongly correlated, except for pairs CBS-CDS, so we drop CDS from the analysis and avoid using them in the same model to avoid multicollinearity issues. These are just comments during the intervention period, which, unsurprisingly are also a good proxy for comments before and comments after.

2 Causal inference

To identify the right variables to condition (or not condition) on to identify the causal effect of the interventions, we first need to think about the causal structure of the problem. Here's a plausible causal structure that we will be working with:

```
dag <- dagitty("
  dag{
    CDS -> ADS -> IC ons
    U [unobserved]
    U -> CBS -> ABS
    U -> ABS
    U -> CDS -> ADS
    U -> ADS
    U -> CAS -> AAS
    U -> AAS
    IC -> AAS
    IC -> CAS
    IT -> CAS
    IT -> AAS
    CBS -> CDS -> CAS
    ABS -> ADS -> AAS
  }")

set.seed(123)
drawdag(dag)
```



ons

Comments during impact attacks during, which trigger interventions. Unmeasured user features cause comments before, which impact attacks before, and also attacks before directly. Comments during (their impact on ADS is already included) impact attacks during during directly and comments after, which impact attacks after and attacks after directly. Intervention count impacts attacks after and comments after. The same directions of impact are included for intervention type. Finally, comments through time are connected causally, and so are attacks.

We already know not to condition on CDS if we condition on CAS or CBS. What else? IT has no backward paths, but IC does. Let's identify all paths from IC to AAS:

```

paths(dag, from = c("IC"), to = "AAS")

## $paths
## [1] "IC -> AAS"
## [2] "IC -> CAS -> AAS"
## [3] "IC -> CAS <- CDS -> ADS -> AAS"
## [4] "IC -> CAS <- CDS -> ADS <- ABS <- CBS <- U -> AAS"
## [5] "IC -> CAS <- CDS -> ADS <- ABS <- U -> AAS"
## [6] "IC -> CAS <- CDS -> ADS <- U -> AAS"
## [7] "IC -> CAS <- CDS <- CBS -> ABS -> ADS -> AAS"
## [8] "IC -> CAS <- CDS <- CBS -> ABS -> ADS <- U -> AAS"
## [9] "IC -> CAS <- CDS <- CBS -> ABS <- U -> AAS"
## [10] "IC -> CAS <- CDS <- CBS -> ABS <- U -> ADS -> AAS"
## [11] "IC -> CAS <- CDS <- CBS <- U -> AAS"
## [12] "IC -> CAS <- CDS <- CBS <- U -> ABS -> ADS -> AAS"
## [13] "IC -> CAS <- CDS <- CBS <- U -> ADS -> AAS"
## [14] "IC -> CAS <- CDS <- U -> AAS"
## [15] "IC -> CAS <- CDS <- U -> ABS -> ADS -> AAS"
## [16] "IC -> CAS <- CDS <- U -> ADS -> AAS"
## [17] "IC -> CAS <- CDS <- U -> CBS -> ABS -> ADS -> AAS"
## [18] "IC -> CAS <- IT -> AAS"
## [19] "IC -> CAS <- U -> AAS"
## [20] "IC -> CAS <- U -> ABS -> ADS -> AAS"
## [21] "IC -> CAS <- U -> ABS <- CBS -> CDS -> ADS -> AAS"
## [22] "IC -> CAS <- U -> ADS -> AAS"
## [23] "IC -> CAS <- U -> CBS -> ABS -> ADS -> AAS"
## [24] "IC -> CAS <- U -> CBS -> CDS -> ADS -> AAS"
## [25] "IC -> CAS <- U -> CDS -> ADS -> AAS"
## [26] "IC -> CAS <- U -> CDS <- CBS -> ABS -> ADS -> AAS"
## [27] "IC <- ADS -> AAS"
## [28] "IC <- ADS <- ABS <- CBS -> CDS -> CAS -> AAS"
## [29] "IC <- ADS <- ABS <- CBS -> CDS -> CAS <- IT -> AAS"

```



```

## [30] "IC <- ADS <- ABS <- CBS -> CDS -> CAS <- U -> AAS"
## [31] "IC <- ADS <- ABS <- CBS -> CDS <- U -> AAS"
## [32] "IC <- ADS <- ABS <- CBS -> CDS <- U -> CAS -> AAS"
## [33] "IC <- ADS <- ABS <- CBS -> CDS <- U -> CAS <- IT -> AAS"
## [34] "IC <- ADS <- ABS <- CBS <- U -> AAS"
## [35] "IC <- ADS <- ABS <- CBS <- U -> CAS -> AAS"
## [36] "IC <- ADS <- ABS <- CBS <- U -> CAS <- IT -> AAS"
## [37] "IC <- ADS <- ABS <- CBS <- U -> CDS -> CAS -> AAS"
## [38] "IC <- ADS <- ABS <- CBS <- U -> CDS -> CAS <- IT -> AAS"
## [39] "IC <- ADS <- ABS <- U -> AAS"
## [40] "IC <- ADS <- ABS <- U -> CAS -> AAS"
## [41] "IC <- ADS <- ABS <- U -> CAS <- IT -> AAS"
## [42] "IC <- ADS <- ABS <- U -> CBS -> CDS -> CAS -> AAS"
## [43] "IC <- ADS <- ABS <- U -> CBS -> CDS -> CAS <- IT -> AAS"
## [44] "IC <- ADS <- ABS <- U -> CDS -> CAS -> AAS"
## [45] "IC <- ADS <- ABS <- U -> CDS -> CAS <- IT -> AAS"
## [46] "IC <- ADS <- CDS -> CAS -> AAS"
## [47] "IC <- ADS <- CDS -> CAS <- IT -> AAS"
## [48] "IC <- ADS <- CDS -> CAS <- U -> AAS"
## [49] "IC <- ADS <- CDS <- CBS -> ABS <- U -> AAS"
## [50] "IC <- ADS <- CDS <- CBS -> ABS <- U -> CAS -> AAS"
## [51] "IC <- ADS <- CDS <- CBS -> ABS <- U -> CAS <- IT -> AAS"
## [52] "IC <- ADS <- CDS <- CBS <- U -> AAS"
## [53] "IC <- ADS <- CDS <- CBS <- U -> CAS -> AAS"
## [54] "IC <- ADS <- CDS <- CBS <- U -> CAS <- IT -> AAS"
## [55] "IC <- ADS <- CDS <- U -> AAS"
## [56] "IC <- ADS <- CDS <- U -> CAS -> AAS"
## [57] "IC <- ADS <- CDS <- U -> CAS <- IT -> AAS"
## [58] "IC <- ADS <- U -> AAS"
## [59] "IC <- ADS <- U -> ABS <- CBS -> CDS -> CAS -> AAS"
## [60] "IC <- ADS <- U -> ABS <- CBS -> CDS -> CAS <- IT -> AAS"
## [61] "IC <- ADS <- U -> CAS -> AAS"
## [62] "IC <- ADS <- U -> CAS <- IT -> AAS"
## [63] "IC <- ADS <- U -> CBS -> CDS -> CAS -> AAS"
## [64] "IC <- ADS <- U -> CBS -> CDS -> CAS <- IT -> AAS"
## [65] "IC <- ADS <- U -> CDS -> CAS -> AAS"
## [66] "IC <- ADS <- U -> CDS -> CAS <- IT -> AAS"
##
## $open
## [1] TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [25] FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE TRUE TRUE FALSE
## [37] TRUE FALSE TRUE TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE FALSE
## [49] FALSE FALSE FALSE TRUE TRUE FALSE TRUE TRUE FALSE TRUE FALSE FALSE
## [61] TRUE FALSE TRUE FALSE TRUE FALSE

```

Crucially, all backdoor paths go through ADS, which then becomes either a fork or a pipe, so all backdoor paths can be closed by conditioning on ADS. Moreover there is only one directed indirect path, it goes through CAS, so we should not condition on it if we are to identify causal effect on attacks mediated by impact on comments (unless we care about the direct effect of IC and IT on AAS, but that's a separate question). This is in line with the adjustment set identified algorithmically, and the same move makes sense when we want to predict CAS.

```

adjustmentSets(dag, exposure = c("IC", "IT"), outcome = "AAS")

## { ADS }

adjustmentSets(dag, exposure = c("IC", "IT"), outcome = "CAS")

## { ADS }

```

It's open season for other variables, and our decision to include them in the model will be guided by information-theoretic criteria of predictive power.

In fact, we will be predicting the difference between attacks before and after, and the difference between comments, before and after. Let's add them to the dag to double-check our selection of variables.

```

dag2 <- dagitty("
  dag{
    CDS -> ADS -> IC   ons
    U [unobserved]
    U -> CBS -> ABS

```

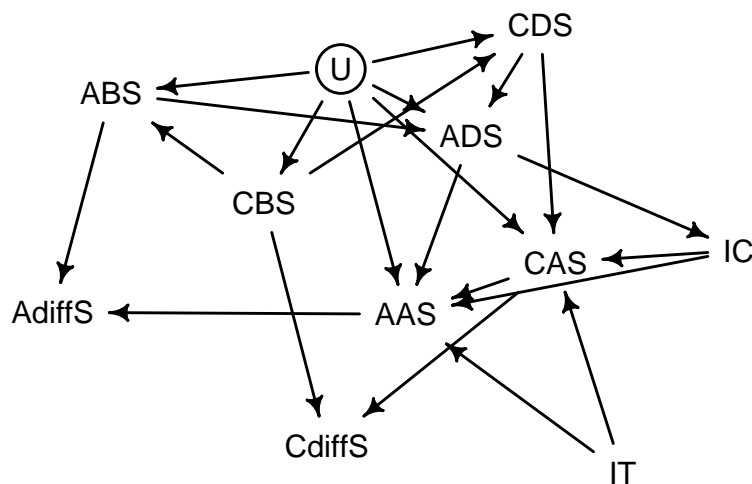
```

U -> ABS
U -> CDS -> ADS
U -> ADS
U -> CAS -> AAS
U -> AAS
IC -> AAS
IC -> CAS
IT -> CAS
IT -> AAS
CBS -> CDS -> CAS
ABS -> ADS -> AAS
ABS -> AdiffS
AAS -> AdiffS
CBS -> CdiffS
CAS -> CdiffS
})

set.seed(123)
drawdag(dag2)
adjustmentSets(dag2, exposure = c("IC", "IT"), outcome = "AdiffS")
## { ADS }

adjustmentSets(dag2, exposure = c("IC", "IT"), outcome = "CdiffS")
## { ADS }

```



ons

3 Bayesian models, priors and diagnostics

We will focus on a class of additive models where the outcome variable is normally distributed around the predicted mean, which is a linear function of predictors (possibly with some interactions). To spoil the story, we will end up using a model, whose specification is as follows:

$$\begin{aligned}
\text{AdiffS} &\sim \text{Norm}(\mu, \sigma) \\
\mu_i &= \alpha + \beta_{\text{ADS}}[\text{group}_i] \times \text{ADS} + \beta_{\text{group}_i} + \beta_{\text{IC}}[\text{group}_i] \times \text{IC} + \\
&\quad + \beta_{\text{ADSIC}} \times \text{ADS} \times \text{IC} + \beta_{\text{CBS}}[\text{group}_i] \times \text{CBS} \\
\alpha &\sim \text{Norm}(0, .3) \\
\beta_{\text{ADS}}[\text{group}_i] &\sim \text{Norm}(0, .3) \\
\beta_{\text{group}_i} &\sim \text{Norm}(0, .3) \\
\beta_{\text{IC}}[\text{group}_i] &\sim \text{Norm}(0, .3) \\
\beta_{\text{ADSIC}} &\sim \text{Norm}(0, .3) \\
\beta_{\text{CBS}}[\text{group}_i] &\sim \text{Norm}(0, .3)
\end{aligned}$$

That is, we take the resulting mean to be the result of the general average (α) and the impact of the following coefficients: group-specific coefficient for ADS, group coefficient, group-specific coefficient for IC, interaction coefficient for ADS and IC, and group-specific coefficient for CBS. This is plausible *prima facie* which group a user belongs to might have impact on how attacks during the treatment is related to attacks after, the role of the intervention count, and the role of comments before. Moreover, the levels of aggressive behavior displayed by the user during treatment might have impact on the role played by the intervention count. Later on we will see that there are information-theoretic reasons to include these interactions.

Now for the priors. One might be suspicious of $\sigma = .3$ we employed and suggest using standard normal distributions with $\sigma = 1$ instead. However, a quick prior predictive check shows that this results in insanely wide priors that are completely unrealistic. (For computational reasons, instead of running the simulations, we load pre-compiled models, but we include the code used to build them).

```
# building model with sd=1
# InteractionsModelDiffSD1 <- ulam(
#   alist(
#     AdiffS ~ dnorm( mu, sigma ),
#     mu <- a + bADS[groupID] * ADS + bIT[groupID] + bIC[groupID] * IC +
#     bADSIC * ADS * IC + bCBS[groupID] * CBS,
#     a ~ dnorm (0,1),
#     bADS[groupID] ~ dnorm(0,1),
#     bADSIC ~ dnorm(0,1),
#     bCBS[groupID] ~ dnorm(0,1),
#     bIT[groupID] ~ dnorm(0,1),
#     bIC[groupID] ~ dnorm(0,1),
#     sigma ~ dexp(1)
#   ),
#   data = summaries
# )
#
# saveRDS(InteractionsModelDiffSD1, file = "models/InteractionsModelDiffSD1.rds")
InteractionsModelDiffSD1 <- readRDS(file = "models/InteractionsModelDiffSD1.rds")

#now model with prior sd = .3
# InteractionsModelDiff <- ulam(
#   alist(
#     AdiffS ~ dnorm( mu, sigma ),
#     mu <- a + bADS[groupID] * ADS + bIT[groupID] + bIC[groupID] * IC +
#     bADSIC * ADS * IC + bCBS[groupID] * CBS,
#     a ~ dnorm (0,0.3),
#     bADS[groupID] ~ dnorm(0,.3),
#     bADSIC ~ dnorm(0,.3),
#     bCBS[groupID] ~ dnorm(0,.3),
#     bIT[groupID] ~ dnorm(0,.3),
#     bIC[groupID] ~ dnorm(0,.3),
#     sigma ~ dexp(1)
#   ),
#   data = summaries
# )
```

```

#saveRDS(InteractionsModelDiff, file = "models/InteractionsModelDiff.rds")

InteractionsModelDiff <- readRDS(file = "models/InteractionsModelDiff.rds")

##prior predictive checks sd =1
ADS <- 0
CBS <- 0
groupID <- 1:3
IC <- 5 #mean for interventions in treatment
data <- expand.grid(ADS = ADS,groupID = groupID, CBS = CBS, IC = IC)
prior <- extract.prior(InteractionsModelDiffSD1, n = 1e4)
mu <- link( InteractionsModelDiffSD1 , post=prior , data=data )
colnames(mu) <- levels(summaries$group)
muLong <- melt(mu)
colnames(muLong) <- c("id", "group", "AdiffS")

priorGroupsSD1 <- ggplot(muLong)+
  geom_violin(aes(x = group, y = AdiffS))+
  theme_tufte()+xlab("")+
  labs(title = "Simulated priors by group",
        subtitle = "(at ADS = CBS = 0, IC at mean = 5, sd = 1)")+
  ylab("change in attacks (standardized)")

ADS <- 0
CBS <- 0
groupID <- 1:3
IC <- 0:20
data <- expand.grid(ADS = ADS,groupID = groupID, CBS = CBS, IC = IC)

prior <- extract.prior(InteractionsModelDiffSD1, n = 1e4)


## recompiling to avoid crashing R session
mu <- link(InteractionsModelDiffSD1 , post=prior , data=data )
mu.mean <- apply( mu , 2, mean )
mu.HPDI <- data.frame(t(apply( mu , 2 , HPDI )))
priorDF <- cbind(data, mu.mean, mu.HPDI)
priorDF$groupID <- as.factor(groupID)
levels(priorDF$groupID) <- c("control", "empathy", "normative")
colnames(priorDF)[2]<- "group"

priorICSD1 <- ggplot(priorDF, aes(x = IC, y = mu.mean, fill = group))+
  geom_line()+geom_ribbon(aes(ymin = X.0.89, ymax = X0.89.), alpha = 0.2)+
  theme_tufte()+ylab("change in attacks (standardized)")+
  labs(title = "Simulated priors for AAS vs IC",
        subtitle = "(at ADS = CBS = 0, sd = 1)")+xlab("interventions")

priorJoint1 <- ggarrange(priorGroupsSD1,priorICSD1, ncol = 2)
priorJoint1Titled <- annotate_figure(priorJoint1,
  top = text_grob("Predictive priors with sd=1 are insanely wide",
    size = 14))
priorJoint1Titled

```

Predictive priors with sd=1 are insanely wide



Some experimentation leads to the value of $\sigma = 3$, which leads to the following priors:

```
#prior predictive check sd =.3
ADS <- 0
CBS <- 0
groupID <- 1:3
IC <- 5 #mean for interventions in treatment
data <- expand.grid(ADS = ADS, groupID = groupID, CBS = CBS, IC = IC)
prior <- extract.prior(InteractionsModelDiff, n = 1e4)
mu <- link(InteractionsModelDiff, post=prior, data=data)
colnames(mu) <- levels(summaries$group)
muLong <- melt(mu)
colnames(muLong) <- c("id", "group", "Adiffs")
head(muLong)

priorGroupSD03 <- ggplot(muLong)+
  geom_violin(aes(x = group, y = Adiffs))+theme_tufte()+
  xlab("")+
  labs(title = "Simulated priors by group",
        subtitle = "(at ADS = CBS = 0, IC at mean = 5, sd = .3)"+
  ylab("change in attacks (standarized)")

ADS <- 0
CBS <- 0
groupID <- 1:3
IC <- 5 #mean for interventions in treatment
data <- expand.grid(ADS = ADS, groupID = groupID, CBS = CBS, IC = IC)
prior <- extract.prior(InteractionsModelDiffSD1, n = 1e4)
mu <- link(InteractionsModelDiffSD1, post=prior, data=data)
colnames(mu) <- levels(summaries$group)
muLong <- melt(mu)
colnames(muLong) <- c("id", "group", "Adiffs")
head(muLong)

priorICSD03 <- ggplot(muLong)+
  geom_violin(aes(x = group, y = Adiffs))+
  theme_tufte()+xlab("")+
  labs(title = "Simulated priors by group",
        subtitle = "(at ADS = CBS = 0, IC at mean = 5, sd = 1)"+
  ylab("change in attacks (standardized)")

priorJoint03 <- ggarrange(priorGroupSD03, priorICSD03, ncol = 2)
```

```
priorJoint03Titled <- annotate_figure(priorJoint03,
  top = text_grob("Predictive priors with sd=.3 seem sensible",
    size = 14))
priorJoint03Titled
```


Predictive priors with sd=.3 seem sensible

Simulated priors by group

(at ADS = CBS = 0, IC at mean = 5, sd = .3)


Simulated priors by group

(at ADS = CBS = 0, IC at mean = 5, sd = 1)



Now, some model diagnostics before we move on. What we are witnessing is (1) stationarity (the chains stay mostly in the most probable regions), (2) good mixing (they explore a range of options in the beginning), and (3) convergence (they stabilize as they progress).


```
traceplot( InteractionsModelDiff )
```



Finally, let's inspect the distribution of residuals. That is, we calculate all predictions, their distance from the actual values, and inspect the distribution of the distances:

```
mu <- link(InteractionsModelDiff)
mu_mean <- apply(mu, 2, mean)
mu_resid <- summaries$AdiffS - mu_mean
ggplot()+geom_density(aes(x = mu_resid))+theme_tufte()+
  ggtitle("Residuals are approximately normally distributed")
```

Residuals are approximately normally distributed



References